

Article Title

Marin Benčević^{1*}, Irena Galić^{1†} and Aleksandra Pizurica^{2†}

¹Faculty of Electrical Engineering, Computer Science and Information Technology, J. J. Strossmayer University of Osijek, Osijek, Croatia.

²TELIN-GAIM, Faculty of Engineering and Architecture, Ghent University, Ghent, Belgium.

*Corresponding author(s). E-mail(s): marin.bencevic@ferit.hr;

Contributing authors: irena.galic@ferit.hr;

aleksandra.pizurica@ugent.be;

[†]These authors contributed equally to this work.

Abstract

The abstract serves both as a general introduction to the topic and as a brief, non-technical summary of the main results and their implications. Authors are advised to check the author instructions for the journal they are submitting to for word limits and if structural elements like subheadings, citations, or equations are permitted.

Keywords: keyword1, Keyword2, Keyword3, Keyword4

1 Motivation

One of the largest problems in medical image processing is the lack of annotated data. To function robustly and to show their generalizability potential, deep learning networks require a large amount of annotated images. However, annotating medical imaging is often time-consuming, precise work. There is a need to improve the data-efficiency and robustness of deep learning networks trained on smaller datasets. Furthermore, there is a need to make the labeling process faster to save experts' time. This paper represents a step towards both of those goals. First, we present a method which uses self-supervised learning to extract salient information from *unlabeled* images which can then be used to

more easily train a deep learning network on a more limited dataset of *labeled* images.

Secondly, we evaluate this on an object detection dataset. In terms of labeling complexity, it is simpler and quicker to label a dataset with bounding boxes than to label each instance as is needed for semantic segmentation. In a lot of medical imaging tasks, a precise semantic segmentation map is not required, and a bounding box communicates sufficient information for further diagnosis, treatment or research on a given image. By improving data efficiency for bounding box labels, we hope to increase their usefulness and thus save time by allowing experts to use bounding box labels in place of semantic segmentation maps.

2 Dataset description and demographics

The dataset used in this paper is a dataset of 15,000 labeled chest radiographs called VinDr-CXR, described in more detail in [?]. While the original dataset contains 3,000 additional test images, we were not able to obtain the labels for these images, and they were not used in this paper. Each scan of the dataset was labeled by three separate radiologists. The dataset was collected from two major Vietnamese hospitals.

We chose this dataset as a good indicator of the generalizability of the findings in this paper due to several reasons. Firstly, to our knowledge, this is the largest radiograph dataset with bounding box labels for each finding. Secondly, having multiple labels for each image allows us to perform a large variety of statistical evaluation methods, such as comparing the inter-observer variability. Having multiple labels also allows us to use automatic methods to find a consensus between them, possibly leading to less noisy annotations.

2.1 Data preparation

Each DICOM image from the dataset was resized to a resolution of 512×512 pixels. We discard all examples for which there is no anomaly found (examples labeled as "no finding"). After discarding, we are left with a total of 4,394 images. We randomly split this dataset into a training set (70%, 3075 images), validation set (10%, 440 images) and a test set (20%, 878 images). The training set was used to train the models, the validation set was used to tune the model hyperparameters and determine when to stop training, and the test set is used for final evaluation. The model did not have access to the test set during training.

The original dataset differentiates 14 different classes, one of which is the "no finding" class. We discard this class, resulting in 13 possible class labels. The class distribution of the full dataset is shown in Figure ??.

Each image can have one or more labels from multiple experts, and these labels often overlap. To produce the least noisy labels, we fused overlapping labels from multiple experts into one label by finding an average rectangle

of several overlapping rectangles. To determine if two rectangles are overlapping, an intersection-over-union (IoU) threshold of 20% is used. A rectangle R_i is defined by its top-left corner (X_{i1}, Y_{i1}) and its bottom-right corner (X_{i2}, Y_{i2}) . An average rectangle \bar{R} defined by coordinates (\bar{X}_1, \bar{Y}_1) and (\bar{X}_2, \bar{Y}_2) is calculated as follows:

$$\begin{aligned}\bar{X}_1 &= \frac{\sum_{i=1}^N X_{i1}}{N}, \bar{Y}_1 = \frac{\sum_{i=1}^N Y_{i1}}{N}, \\ \bar{X}_2 &= \frac{\sum_{i=1}^N X_{i2}}{N}, \bar{Y}_2 = \frac{\sum_{i=1}^N Y_{i2}}{N}\end{aligned}\quad (1)$$

Where N is the number of rectangles to average. This approach is based on weighted boxes fusion described in [?], but modified such that each bounding box has equal weight and confidence since they were manually labeled by an expert. An example of fused bounding boxes is shown in ??.

3 Methods

The main goal of this paper is to analyze how self-supervised model pre-training affects data efficiency for object detection in medical images. Therefore, we first train a baseline deep learning model with no pre-training and on the full labeled training dataset following a standard approach for this type of problem. This model will be used as a point of comparison to more objectively evaluate the pre-trained models.

To evaluate the pre-training, we randomly split the training dataset into two separate datasets, a pre-training and fine-tuning dataset. For the pre-training dataset we discard all class labels, as this dataset will be used to pre-train the model using self-supervised learning on unlabeled data. The fine-tuning dataset will then be used to fine-tune the pre-trained model using standard supervised learning. We train nine different pre-trained models in total, ranging from 10% to 90% of the total training set in the unlabeled pre-train dataset, in increments of 10%. A summary of our approach is shown in 1.

3.1 Balancing the dataset

As shown in Figure ?? the dataset is highly imbalanced. To improve class-balance during training we oversample examples with less-represented classes. We use the oversampling approach described in [?]. The general approach is as follows:

1. For each class c , compute its frequency $f(c)$ in the training set.
2. Compute a class repeat factor $r(c) = \max(1, \sqrt{t/f(c)})$, where t is a threshold value empirically set to 0.4 in our experiments. The value of 0.4 produced the maximum validation mean average precision for the baseline model.

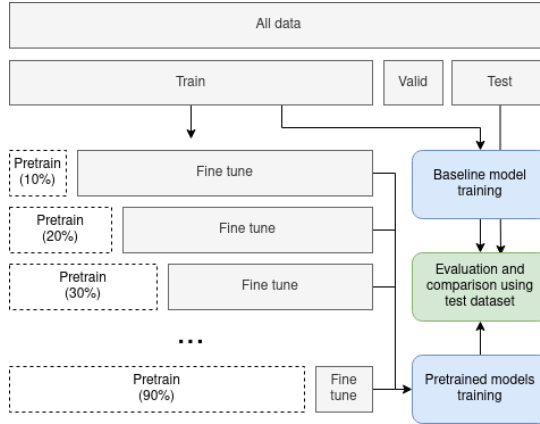


Fig. 1 A summary of our experiments. A percentage of the training dataset is moved to the pre-training dataset and models are pre-trained using the pre-training datasets and then fine tuned with the rest of the training data. The pre-training datasets are unlabeled. A separate baseline model is trained using the full labeled dataset.

- For each image I , compute an image repeat factor $r(I) = C * r(c)$ where $C = \max(r(c_i))$ for each class c_i in image I .

This balancing is performed on-line during training and only on the training set. We found that balancing the dataset in this way drastically improved the mean average precision when averaged across all classes.

3.2 Baseline model details

For an objective and fair comparison, we train a baseline model using a standard deep learning-based approach for object detection. We use a Faster R-CNN-based model [?] with a ResNet50 encoder [?]. The model is initialized using pre-trained weights trained on the COCO dataset for 12 epochs, a batch size of 2 and using the SGD optimizer with a learning rate of 0.0002, momentum of 0.9 and a weight decay of 0.0001.

3.3 Pretraining model details

The pretraining model we use is a SimCLR-based model [?] to pre-train a ResNet-50 backbone, the same backbone used in the baseline model. We train the SimCLR model (described later in this section) using the unlabeled pre-training dataset. We then use the pre-trained backbone in the same model as our baseline model and fine-tune the final model on the fine-tuning dataset. The result is a model similar to our backbone model but trained on fewer labeled data.

SimCLR uses contrastive learning where an example image is augmented randomly twice and each augmentation is fed into a separate encoder branch, where the branches use shared weights. The network outputs two feature maps, one for each augmentation. The loss function measures the difference between

these two feature maps. The closer the two feature maps are, the lower the loss. This ensures that two augmentations from the same example will produce similar feature maps, thus making the network learn salient features which are invariant to the chosen augmentations.

In our experiments, we use the following augmentations for SimCLR training:

1. A random crop and resize of the original image by a factor of 0.2 to 1.
2. A random horizontal flip (with a 50% chance).
3. A random Gaussian blur with σ between 0.1 and 2, and a kernel size of 21.
4. A random amount of Gaussian noise with σ being a random number between 12.5% and 25% of the mean image pixel value.

In addition, each training and validation image has histogram normalization applied.

4 Results

5 Discussion

Discussions should be brief and focused. In some disciplines use of Discussion or ‘Conclusion’ is interchangeable. It is not mandatory to use both. Some journals prefer a section ‘Results and Discussion’ followed by a section ‘Conclusion’. Please refer to Journal-level guidance for any specific requirements.

6 Conclusion

Conclusions may be used to restate your hypothesis or research question, restate your major findings, explain the relevance and the added value of your work, highlight any limitations of your study, describe future directions for research and recommendations.

In some disciplines use of Discussion or ‘Conclusion’ is interchangeable. It is not mandatory to use both. Please refer to Journal-level guidance for any specific requirements.

Supplementary information. If your article has accompanying supplementary file/s please state so here.

Authors reporting data from electrophoretic gels and blots should supply the full unprocessed scans for key as part of their Supplementary information. This may be requested by the editorial team/s if it is missing.

Please refer to Journal-level guidance for any specific requirements.

Acknowledgments. Acknowledgments are not compulsory. Where included they should be brief. Grant or contribution numbers may be acknowledged.

Please refer to Journal-level guidance for any specific requirements.

Declarations

Some journals require declarations to be submitted in a standardised format. Please check the Instructions for Authors of the journal to which you are submitting to see if you need to complete this section. If yes, your manuscript must contain the following sections under the heading ‘Declarations’:

- Funding
- Conflict of interest/Competing interests (check journal-specific guidelines for which heading to use)
- Ethics approval
- Consent to participate
- Consent for publication
- Availability of data and materials
- Code availability
- Authors’ contributions

If any of the sections are not relevant to your manuscript, please include the heading and write ‘Not applicable’ for that section.

7

8 Equations

Equations in L^AT_EX can either be inline or on-a-line by itself (“display equations”). For inline equations use the $\$ \dots \$$ commands. E.g.: The equation $H\psi = E\psi$ is written via the command `$\$H \ \psi = E \ \psi\$$` .

For display equations (with auto generated equation numbers) one can use the `equation` or `align` environments:

$$\|\tilde{X}(k)\|^2 \leq \frac{\sum_{i=1}^p \|\tilde{Y}_i(k)\|^2 + \sum_{j=1}^q \|\tilde{Z}_j(k)\|^2}{p+q}. \quad (2)$$

where,

$$\begin{aligned} D_\mu &= \partial_\mu - ig \frac{\lambda^a}{2} A_\mu^a \\ F_{\mu\nu}^a &= \partial_\mu A_\nu^a - \partial_\nu A_\mu^a + gf^{abc} A_\mu^b A_\nu^a \end{aligned} \quad (3)$$

Notice the use of `\nonumber` in the `align` environment at the end of each line, except the last, so as not to produce equation numbers on lines where no equation numbers are required. The `\label{}` command should only be used at the last line of an `align` environment where `\nonumber` is not used.

$$Y_\infty = \left(\frac{m}{\text{GeV}}\right)^{-3} \left[1 + \frac{3 \ln(m/\text{GeV})}{15} + \frac{\ln(c_2/5)}{15}\right] \quad (4)$$

The class file also supports the use of `\mathbb{}`, `\mathscr{}` and `\mathcal{}` commands. As such `\mathbb{R}`, `\mathscr{R}` and `\mathcal{R}` produces \mathbb{R} , \mathscr{R} and \mathcal{R} respectively (refer Subsubsection ??).

9 Tables

Tables can be inserted via the normal table and tabular environment. To put footnotes inside tables you should use `\footnotetext[]{\dots}` tag. The footnote appears just below the table itself (refer Tables 1 and 2). For the corresponding footnotemark use `\footnotemark[...]`

Table 1 Caption text

Column 1	Column 2	Column 3	Column 4
row 1	data 1	data 2	data 3
row 2	data 4	data 5 ¹	data 6
row 3	data 7	data 8	data 9 ²

Source: This is an example of table footnote.
This is an example of table footnote.

¹Example for a first table footnote. This is an example of table footnote.

²Example for a second table footnote. This is an example of table footnote.

The input format for the above table is as follows:

```
\begin{table}[<placement-specifier>]
\begin{center}
\begin{minipage}{<preferred-table-width>}
\caption{<table-caption>}\label{<table-label>}%
\begin{tabular}{@{}llll@{}}
\toprule
Column 1 & Column 2 & Column 3 & Column 4\\
\midrule
row 1 & data 1 & data 2 & data 3 \\
row 2 & data 4 & data 5\footnotemark[1] & data 6 \\
row 3 & data 7 & data 8 & data 9\footnotemark[2]\\
\botrule
\end{tabular}
\footnotetext{Source: This is an example of table footnote.
This is an example of table footnote.}
\footnotetext[1]{Example for a first table footnote.
This is an example of table footnote.}
\footnotetext[2]{Example for a second table footnote.
This is an example of table footnote.}
```

```
\end{minipage}
\end{center}
\end{table}
```

Table 2 Example of a lengthy table which is set to full textwidth

Project	Element 1 ¹			Element 2 ²		
	Energy	σ_{calc}	σ_{expt}	Energy	σ_{calc}	σ_{expt}
Element 3	990 A	1168	1547 ± 12	780 A	1166	1239 ± 100
Element 4	500 A	961	922 ± 10	900 A	1268	1092 ± 40

Note: This is an example of table footnote. This is an example of table footnote this is an example of table footnote this is an example of table footnote.

¹Example for a first table footnote.

²Example for a second table footnote.

In case of double column layout, tables which do not fit in single column width should be set to full text width. For this, you need to use `\begin{table*}` ... `\end{table*}` instead of `\begin{table}` ... `\end{table}` environment. Lengthy tables which do not fit in textwidth should be set as rotated table. For this, you need to use `\begin{sidewaystable}` ... `\end{sidewaystable}` instead of `\begin{table*}` ... `\end{table*}` environment. This environment puts tables rotated to single column width. For tables rotated to double column width, use `\begin{sidewaystable*}` ... `\end{sidewaystable*}`.

10 Figures

As per the L^AT_EX standards you need to use eps images for L^AT_EX compilation and pdf/jpg/png images for PDFLaTeX compilation. This is one of the major difference between L^AT_EX and PDFLaTeX. Each image should be from a single input .eps/vector image file. Avoid using subfigures. The command for inserting images for L^AT_EX and PDFLaTeX can be generalized. The package used to insert images in LaTeX/PDFLaTeX is the graphicx package. Figures can be inserted via the normal figure environment as shown in the below example:

```
\begin{figure}[<placement-specifier>]
\centering
\includegraphics{<eps-file>}
\caption{<figure-caption>}\label{<figure-label>}
\end{figure}
```


Table 3 Tables which are too long to fit, should be written using the “sidewaystable” environment as shown here

Projectile	Element 1 ¹		Element ²	
	Energy	σ_{calc}	Energy	σ_{expt}
Element 3	990 A	1168	780 A	1239 ± 100
Element 4	500 A	961	900 A	1092 ± 40
Element 5	990 A	1168	780 A	1239 ± 100
Element 6	500 A	961	900 A	1092 ± 40

Note: This is an example of table footnote this is an example of table footnote this is an example of table footnote this is an example of table footnote this is an example of table footnote.

¹This is an example of table footnote.



Fig. 2 This is a widefig. This is an example of long caption this is an example of long caption this is an example of long caption this is an example of long caption

In case of double column layout, the above format puts figure caption-s/images to single column width. To get spanned images, we need to provide `\begin{figure*} ... \end{figure*}`.

For sample purpose, we have included the width of images in the optional argument of `\includegraphics` tag. Please ignore this.

11 Algorithms, Program codes and Listings

Packages `algorithm`, `algorithmicx` and `algpseudocode` are used for setting algorithms in L^AT_EX using the format:

```
\begin{algorithm}
\caption{<alg-caption>}\label{<alg-label>}
\begin{algorithmic}[1]
. . .
\end{algorithmic}
\end{algorithm}
```

You may refer above listed package documentations for more details before setting `algorithm` environment. For program codes, the “program” package is required and the command to be used is `\begin{program} ... \end{program}`. A fast exponentiation procedure:

```
begin
  for  $i := 1$  to 10 step 1 do
     $\text{expt}(2, i)$ ;
     $\text{newline}()$  od           Comments will be set flush to the right margin
where
proc  $\text{expt}(x, n)$   $\equiv$ 
   $z := 1$ ;
  do if  $n = 0$  then exit fi;
  do if  $\text{odd}(n)$  then exit fi;
    comment: This is a comment statement;
     $n := n/2$ ;  $x := x * x$  od;
   $\{n > 0\}$ ;
   $n := n - 1$ ;  $z := z * x$  od;
  print( $z$ ).
end
```

Algorithm 1 Calculate $y = x^n$

Require: $n \geq 0 \vee x \neq 0$ **Ensure:** $y = x^n$

```

1:  $y \leftarrow 1$ 
2: if  $n < 0$  then
3:    $X \leftarrow 1/x$ 
4:    $N \leftarrow -n$ 
5: else
6:    $X \leftarrow x$ 
7:    $N \leftarrow n$ 
8: end if
9: while  $N \neq 0$  do
10:  if  $N$  is even then
11:     $X \leftarrow X \times X$ 
12:     $N \leftarrow N/2$ 
13:  else [ $N$  is odd]
14:     $y \leftarrow y \times X$ 
15:     $N \leftarrow N - 1$ 
16:  end if
17: end while

```

Similarly, for listings, use the listings package. `\begin{lstlisting}` ... `\end{lstlisting}` is used to set environments similar to `verbatim` environment. Refer to the `lstlisting` package documentation for more details.

```

for i:=maxint to 0 do
begin
{ do nothing }
end;
Write( 'Case_insensitive.' );
Write( 'Pascal_keywords.' );

```

12 Cross referencing

Environments such as figure, table, equation and align can have a label declared via the `\label{#label}` command. For figures and table environments use the `\label{}` command inside or just below the `\caption{}` command. You can then use the `\ref{#label}` command to cross-reference them. As an example, consider the label declared for Figure 2 which is `\label{fig1}`. To cross-reference it, use the command `Figure \ref{fig1}`, for which it comes up as “Figure 2”.

Example 1 Phasellus adipiscing semper elit. Proin fermentum massa ac quam. Sed diam turpis, molestie vitae, placerat a, molestie nec, leo. Maecenas lacinia. Nam ipsum ligula, eleifend at, accumsan nec, suscipit a, ipsum. Morbi blandit ligula feugiat magna. Nunc eleifend consequat lorem.

Sample body text. Sample body text. Sample body text. Sample body text. Sample body text. Sample body text. Sample body text. Sample body text.

Remark 1 Phasellus adipiscing semper elit. Proin fermentum massa ac quam. Sed diam turpis, molestie vitae, placerat a, molestie nec, leo. Maecenas lacinia. Nam ipsum ligula, eleifend at, accumsan nec, suscipit a, ipsum. Morbi blandit ligula feugiat magna. Nunc eleifend consequat lorem.

Sample body text. Sample body text. Sample body text. Sample body text. Sample body text. Sample body text. Sample body text. Sample body text.

Definition 1 (Definition sub head) Example definition text. Example definition text. Example definition text. Example definition text. Example definition text. Example definition text. Example definition text.

Additionally a predefined “proof” environment is available: `\begin{proof} ... \end{proof}`. This prints a “Proof” head in italic font style and the “body text” in roman font style with an open square at the end of each proof environment.

Proof Example for proof text. Example for proof text. Example for proof text. Example for proof text. Example for proof text. Example for proof text. Example for proof text. Example for proof text. Example for proof text. Example for proof text. □

Sample body text. Sample body text. Sample body text. Sample body text. Sample body text. Sample body text. Sample body text. Sample body text.

Proof of Theorem 1 Example for proof text. Example for proof text. Example for proof text. Example for proof text. Example for proof text. Example for proof text. Example for proof text. Example for proof text. Example for proof text. Example for proof text. □

For a quote environment, use `\begin{quote} ... \end{quote}`

Quoted text example. Aliquam porttitor quam a lacus. Praesent vel arcu ut tortor cursus volutpat. In vitae pede quis diam bibendum placerat. Fusce elementum convallis neque. Sed dolor orci, scelerisque ac, dapibus nec, ultricies ut, mi. Duis nec dui quis leo sagittis commodo.

Sample body text. Sample body text. Sample body text. Sample body text. Sample body text. Sample body text (refer Figure 2). Sample body text. Sample body text. Sample body text (refer Table 3).

14 Methods

Topical subheadings are allowed. Authors must ensure that their Methods section includes adequate experimental and characterization data necessary for others in the field to reproduce their work. Authors are encouraged to include RIIIDs where appropriate.

Ethical approval declarations (only required where applicable) Any article reporting experiment/s carried out on (i) live vertebrate (or higher invertebrates), (ii) humans or (iii) human samples must include an unambiguous statement within the methods section that meets the following requirements:

1. Approval: a statement which confirms that all experimental protocols were approved by a named institutional and/or licensing committee. Please identify the approving body in the methods section
2. Accordance: a statement explicitly saying that the methods were carried out in accordance with the relevant guidelines and regulations
3. Informed consent (for experiments involving humans or human tissue samples): include a statement confirming that informed consent was obtained from all participants and/or their legal guardian/s

If your manuscript includes potentially identifying patient/participant information, or if it describes human transplantation research, or if it reports results of a clinical trial then additional information will be required. Please visit (<https://www.nature.com/nature-research/editorial-policies>) for Nature Portfolio journals, (<https://www.springer.com/gp/authors-editors/journal-author/journal-author-helpdesk/publishing-ethics/14214>) for Springer Nature journals, or (<https://www.biomedcentral.com/getpublished/editorial-policies#ethics+and+consent>) for BMC.

15 Discussion

Discussions should be brief and focused. In some disciplines use of Discussion or ‘Conclusion’ is interchangeable. It is not mandatory to use both. Some journals prefer a section ‘Results and Discussion’ followed by a section ‘Conclusion’. Please refer to Journal-level guidance for any specific requirements.

16 Conclusion

Conclusions may be used to restate your hypothesis or research question, restate your major findings, explain the relevance and the added value of your work, highlight any limitations of your study, describe future directions for research and recommendations.

In some disciplines use of Discussion or ‘Conclusion’ is interchangeable. It is not mandatory to use both. Please refer to Journal-level guidance for any specific requirements.

Supplementary information. If your article has accompanying supplementary file/s please state so here.

Authors reporting data from electrophoretic gels and blots should supply the full unprocessed scans for key as part of their Supplementary information. This may be requested by the editorial team/s if it is missing.

Please refer to Journal-level guidance for any specific requirements.

Acknowledgments. Acknowledgments are not compulsory. Where included they should be brief. Grant or contribution numbers may be acknowledged.

Please refer to Journal-level guidance for any specific requirements.

Declarations

Some journals require declarations to be submitted in a standardised format. Please check the Instructions for Authors of the journal to which you are submitting to see if you need to complete this section. If yes, your manuscript must contain the following sections under the heading ‘Declarations’:

- Funding
- Conflict of interest/Competing interests (check journal-specific guidelines for which heading to use)
- Ethics approval
- Consent to participate
- Consent for publication
- Availability of data and materials
- Code availability
- Authors’ contributions

If any of the sections are not relevant to your manuscript, please include the heading and write ‘Not applicable’ for that section.

Editorial Policies for:

Springer journals and proceedings:

<https://www.springer.com/gp/editorial-policies>

Nature Portfolio journals:

<https://www.nature.com/nature-research/editorial-policies>

Scientific Reports:

<https://www.nature.com/srep/journal-policies/editorial-policies>

BMC journals:

<https://www.biomedcentral.com/getpublished/editorial-policies>

Appendix A Section title of first appendix

An appendix contains supplementary information that is not an essential part of the text itself but which may be helpful in providing a more comprehensive understanding of the research problem or it is information that is too cumbersome to be included in the body of the paper.